

Multi-AS Network Simulation in ns-3: BGP-like Routing, Link Failure Recovery, and Performance Analysis

Salih Sefer — 820230313

May 2026

Contents

1	Abstract	3
2	Topology Description	4
2.1	Autonomous System Overview	4
2.2	IP Addressing Plan	5
3	Routing Structure	6
3.1	Intra-AS: OSPF-like Shortest-Path Routing	6
3.2	Inter-AS: BGP-like Policy Routing	6
3.3	Link Failure and Recovery Workflow	7
4	Experimental Setup	9
4.1	Scenario Matrix	9
4.2	Traffic Generation	9
4.3	Failure and Recovery Timeline	9
4.4	Simulation Parameters	10
5	Results	11
5.1	Summary Statistics	11
5.2	Performance Plots	11
6	Analysis and Interpretation	15
6.1	Effect of Node Count on Network Performance	15
6.2	Balanced vs. Unbalanced Distribution Impact	15
6.3	Topology Change Response: Convergence Analysis	16
6.4	Observable Bottlenecks: AS2 as Transit Hub	16
7	Reproducibility	17
7.1	Environment Setup	17
7.2	Running the Simulations	17
7.3	Generating Analysis Outputs	17
7.4	Visualising with NetAnim	18
8	Conclusion	19

9	Appendix A: Repository Directory Structure	20
10	Appendix B: Command-Line Flags	20
11	Appendix C: RNG Seed and Run Configuration	21

1 Abstract

This report presents a parametric multi-autonomous-system (Multi-AS) network simulation built on ns-3 42. Three interconnected autonomous systems are modelled with randomised intra-AS topologies, BGP-like inter-AS policy routing, and a timed link-failure/recovery scenario. Six scenarios are evaluated across two node-count distribution strategies (balanced and unbalanced) at three scale points (20, 50, and 100 nodes), each replicated three times with distinct RNG seeds. Per-flow metrics — end-to-end delay, throughput, and packet loss — together with the BGP convergence time following a primary-link failure are collected via ns-3 FlowMonitor and aggregated into a comparative analysis. Results show that network scale and distribution asymmetry have measurable but bounded effects on delay and throughput, while convergence time is determined primarily by the simulated BGP hold-timer delay of 2–5 s.

2 Topology Description

2.1 Autonomous System Overview

The simulation models three ASes interconnected by five point-to-point links. Each AS contains two border routers (BR_a and BR_b at node indices 0 and 1 within the AS) plus a variable number of internal routers built using a randomised two-group spanning tree that guarantees two vertex-disjoint paths between every BR pair.

```
graph LR
  subgraph AS1["AS 1 · 10.1.0.0/16"]
    direction TB
    A1b["BR_b (idx 1)"]
    A1a["BR_a (idx 0)"]
    A1i(["internal routers\n×(N_1-2)"])
  end
  subgraph AS2["AS 2 · 10.2.0.0/16"]
    direction TB
    A2b["BR_b (idx 1)"]
    A2a["BR_a (idx 0)"]
    A2i(["internal routers\n×(N_2-2)"])
  end
  subgraph AS3["AS 3 · 10.3.0.0/16"]
    direction TB
    A3b["BR_b (idx 1)"]
    A3a["BR_a (idx 0)"]
    A3i(["internal routers\n×(N_3-2)"])
  end

  A1a -- "1 Gbps / 5 ms · metric 50 [primary]" --- A2a
  A1b -- "100 Mbps / 15 ms · metric 50 [backup]" --- A2b
  A2a -- "1 Gbps / 5 ms · metric 50 [primary]" --- A3a
  A2b -- "100 Mbps / 15 ms · metric 50 [backup]" --- A3b
  A1a -. "1 Gbps / 5 ms · metric 30 [DIRECT — fails at t=20 s]" .-> A3a
```

Figure 1. Inter-AS topology. Dashed edge (AS1.BR_a ↔ AS3.BR_a) is the direct link that fails at t = 20 s; traffic then reroutes via AS2 (metric 80 backup). Render with [Mermaid Live](#) or VS Code *Markdown Preview Enhanced* for a graphical view.

The five inter-AS links are ordered and referred to by index throughout the codebase:

Idx	Endpoints	Bandwidth	Delay	Role
0	AS1.BR_a ↔ AS2.BR_a	1 Gbps	5 ms	AS1–AS2 primary
1	AS1.BR_b ↔ AS2.BR_b	100 Mbps	15 ms	AS1–AS2 backup
2	AS2.BR_a ↔ AS3.BR_a	1 Gbps	5 ms	AS2–AS3 primary
3	AS2.BR_b ↔ AS3.BR_b	100 Mbps	15 ms	AS2–AS3 backup
4	AS1.BR_a ↔ AS3.BR_a	1 Gbps	5 ms	AS1–AS3 direct (fails)

Table 1. Inter-AS link parameters. Intra-AS links use 100 Mbps / 2 ms uniformly.

2.2 IP Addressing Plan

Each AS owns a private /16 block for its intra-AS /30 subnets. The five inter-AS point-to-point links share a sequential 172.16.1.0/24 address space, advancing by 4 for each new link.

Block	Assignment	Subnet per link
10.1.0.0/16	AS1 intra-AS links	/30, allocated by <code>NewNetwork()</code>
10.2.0.0/16	AS2 intra-AS links	/30, allocated by <code>NewNetwork()</code>
10.3.0.0/16	AS3 intra-AS links	/30, allocated by <code>NewNetwork()</code>
172.16.1.0/24+	Inter-AS links (all)	/30, sequential

Table 2. IP addressing plan.

Specific inter-AS host addresses (first/second usable address in each /30):

Link	Subnet	AS-left IP	AS-right IP
0	172.16.1.0/30	172.16.1.1 (AS1.BR_a)	172.16.1.2 (AS2.BR_a)
1	172.16.1.4/30	172.16.1.5 (AS1.BR_b)	172.16.1.6 (AS2.BR_b)
2	172.16.1.8/30	172.16.1.9 (AS2.BR_a)	172.16.1.10 (AS3.BR_a)
3	172.16.1.12/30	172.16.1.13 (AS2.BR_b)	172.16.1.14 (AS3.BR_b)
4	172.16.1.16/30	172.16.1.17 (AS1.BR_a)	172.16.1.18 (AS3.BR_a)

Table 3. Inter-AS /30 host addresses.

3 Routing Structure

3.1 Intra-AS: OSPF-like Shortest-Path Routing

All nodes within each AS run ns-3's `Ipv4GlobalRouting`, which computes Dijkstra-based shortest-path forwarding tables over the entire simulated topology. This is installed via `Ipv4GlobalRoutingHelper::PopulateRoutingTables()` after the full node/link graph is built, analogous to an OSPF domain-wide link-state database synchronisation and SPF computation.

`Ipv4ListRouting` is installed on every node with a priority stack:

```
Priority 10 → Ipv4StaticRouting    (border-router policy routes only)
Priority  5 → Ipv4GlobalRouting    (shortest-path for everything else)
```

For non-BR nodes, the static routing table is empty; all lookups fall through to global routing. For BR nodes, the static routes covering remote AS prefixes (/16 aggregates) are matched first; intra-AS destinations fall through to global routing.

Semantic note. `Ipv4GlobalRouting` is not OSPF: it does not exchange Link-State Advertisements, does not use configurable per-link costs beyond a unit default, and the SPF runs centrally in simulation time rather than in a distributed protocol. The analogy is structural (link-state topology, loop-free forwarding) rather than protocol-level.

3.2 Inter-AS: BGP-like Policy Routing

Six border routers carry static /16 policy routes installed at $t = 0$. The routing intent is modelled after BGP's path-preference mechanism via explicit Dijkstra-metric values: **lower metric = preferred path**.

BR	Destination	Next-hop	Via link	Metric	Role
AS1.BR_a	0.2.0.0/16	AS2.BR_a	link[0]	50	to AS2 primary
AS1.BR_a	0.3.0.0/16	AS3.BR_a	link[4]	30	to AS3 direct ← fails
AS1.BR_a	0.3.0.0/16	AS2.BR_a	link[0]	80	to AS3 via AS2 (backup)
AS1.BR_b	0.2.0.0/16	AS2.BR_b	link[1]	50	to AS2 via backup
AS1.BR_b	0.3.0.0/16	AS2.BR_b	link[1]	80	to AS3 via AS2 backup
AS2.BR_a	0.1.0.0/16	AS1.BR_a	link[0]	50	return to AS1
AS2.BR_a	0.3.0.0/16	AS3.BR_a	link[2]	50	to AS3 primary
AS2.BR_b	0.1.0.0/16	AS1.BR_b	link[1]	50	return to AS1

BR	Destination	Next-hop	Via link	Metric	Role
AS2.BR_b	0.3.0.0/16	AS3.BR_b	link[3]	50	to AS3 backup
AS3.BR_a	0.2.0.0/16	AS2.BR_a	link[2]	50	to AS2 primary
AS3.BR_a	0.1.0.0/16	AS1.BR_a	link[4]	30	to AS1 direct ← fails
AS3.BR_a	0.1.0.0/16	AS2.BR_a	link[2]	80	to AS1 via AS2 (backup)
AS3.BR_b	0.2.0.0/16	AS2.BR_b	link[3]	50	to AS2 backup
AS3.BR_b	0.1.0.0/16	AS2.BR_b	link[3]	80	to AS1 via AS2 backup

Table 4. Border-router policy routes. Ipv4StaticRouting selects the lowest-metric matching entry; higher-metric entries serve as pre-installed backups.

Semantic note. This is not full BGP: there is no TCP session, no OPEN/UPDATE/KEEPALIVE exchange, no route-reflector hierarchy, and no AS_PATH attribute. The simulation captures BGP’s policy-routing semantics (preferred vs. backup paths, hold-timer delay before route withdrawal) while staying within ns-3’s routing API.

3.3 Link Failure and Recovery Workflow

The failure sequence for the AS1_AS3_PRIMARY link is:

```

t = failureTime (default 20 s)
→ BgpLikeRouter::TriggerFailure()
    1. Ipv4::SetDown() on both endpoints of link[4]
       (AS1.BR_a interface and AS3.BR_a interface go DOWN)
    2. MetricsCollector::MarkFailureEvent() – convergence timer starts
    3. Schedule RecomputeBgpTables() after Uniform[2, 5] s delay

t = failureTime + Uniform[2, 5] s
→ BgpLikeRouter::RecomputeBgpTables()
    1. Remove all static routes whose next-hop is reachable only via link[4]
       (metric=30 routes on AS1.BR_a and AS3.BR_a withdrawn)
    2. Ipv4GlobalRoutingHelper::RecomputeRoutingTables()
       (global SPF updated for non-BR nodes)
    3. MetricsCollector::MarkConvergenceEvent() – convergence time recorded
→ Surviving metric=80 routes on AS1.BR_a and AS3.BR_a now best match
→ Traffic path: AS1.BR_a → link[0] → AS2.BR_a → link[2] → AS3.BR_a

t = failureTime + 30 s (= 50 s by default)
→ BgpLikeRouter::TriggerRecovery()

```

1. `Ipv4::SetUp()` restores `link[4]`
2. `Metric=30` routes re-added to BR static tables
3. `RecomputeRoutingTables()` – direct path restored

The BGP hold-timer analogue (`Uniform[2, 5]` s) simulates the real-world delay between link failure detection and route withdrawal propagation in a BGP-speaking network.

4 Experimental Setup

4.1 Scenario Matrix

Node counts are distributed across the three ASes according to the selected strategy.

Scenario	N	Distribution	N_1 (AS1)	N_2 (AS2)	N_3 (AS3)	Flows
S1	20	balanced	7	7	6	6
S2	20	unbalanced	4	6	10	6
S3	50	balanced	17	17	16	9
S4	50	unbalanced	10	15	25	9
S5	100	balanced	34	34	32	14
S6	100	unbalanced	20	30	50	14

Table 5. Scenario matrix. N_i computed by: balanced = $\text{ceil}(N/3) / \text{ceil}(N/3) / \text{remainder}$; unbalanced = $\text{round}(N \times 0.2) / \text{round}(N \times 0.3) / \text{remainder}$. Flows = $4 + \text{floor}(N/10)$.

Each scenario is executed three times (run IDs 1–3) with seeds `BASE_SEED + run_id` (default `BASE_SEED = 1`), yielding 18 independent simulation runs in total.

4.2 Traffic Generation

UDP traffic is modelled with ns-3 `OnOffHelper` (constant-on) and `PacketSinkHelper` pairs. Sources and sinks are placed on non-BR nodes ($\text{index} \geq 2$ within each AS) chosen uniformly at random; BRs are used only if an AS has fewer than three nodes.

Flow set	Directions	Count
Base (always)	AS1→AS2, AS2→AS3, AS1→AS3, AS3→AS1	4
Extra (cycling)	AS2→AS1, AS3→AS2, then repeating base dirs	$\text{floor}(N/10)$

Table 6. Traffic matrix. Each flow: 5 Mbps UDP, 1500 B packet size. Sinks start at $t = 1$ s; sources at $t = 2$ s; all stop at $t = \text{simTime} - 1$ s.

4.3 Failure and Recovery Timeline

```

t = 0 s      Topology built, routing tables populated
t = 2 s      Traffic flows begin (UDP always-on)
t = 20 s     AS1↔AS3 direct link (link[4]) injected as DOWN
              BGP hold-timer starts: Uniform[2, 5] s
t ≈ 22–25 s  BGP convergence completes:
              - Metric-30 direct routes withdrawn from AS1.BR_a, AS3.BR_a
                - Metric-80 backup routes via AS2 become active
              - Traffic path: AS1 → AS2 → AS3 (via link[0] + link[2])
t = 50 s     Link recovery: link[4] restored, metric-
30 routes reinstated

```

```

t = 59 s      All traffic flows stop
t = 60 s      Simulator::Stop(); FlowMonitor statistics serialised

```

4.4 Simulation Parameters

Parameter	Value
Simulation duration (<code>simTime</code>)	60 s
Failure injection time (<code>failureTime</code>)	20 s
Recovery offset	30 s after failure (= 50 s)
BGP convergence delay	Uniform[2, 5] s
Runs per scenario	3
Base RNG seed	1
FlowMonitor histogram/probe output	enabled
NetAnim trace max packets	500 000

Table 7. Fixed simulation parameters.

5 Results

5.1 Summary Statistics

The table below is generated by `scripts/aggregate_results.py` from the 18 simulation runs. Values are mean $\pm$ standard deviation across the three repeated runs of each scenario. | Nodes | Distribution

Delay (ms)	Throughput (Mbps)	Loss (%)	Convergence (s)	Nodes	Distribution
16.984 $\pm$ 2.952	5.092 $\pm$ 0.000	0.002 $\pm$ 0.001	2.782 $\pm$ 0.849	20	balanced
16.931 $\pm$ 1.646	4.980 $\pm$ 0.194	0.003 $\pm$ 0.000	2.782 $\pm$ 0.849	20	unbalanced
23.625 $\pm$ 0.742	5.091 $\pm$ 0.000	0.000 $\pm$ 0.001	3.430 $\pm$ 1.377	50	balanced
17.391 $\pm$ 1.798	5.092 $\pm$ 0.000	0.002 $\pm$ 0.001	3.430 $\pm$ 1.377	50	unbalanced
20.474 $\pm$ 3.136	5.039 $\pm$ 0.092	0.002 $\pm$ 0.000	3.128 $\pm$ 0.852	100	balanced
23.649 $\pm$ 2.119	4.984 $\pm$ 0.093	0.002 $\pm$ 0.000	3.128 $\pm$ 0.852	100	unbalanced

Table 8. Scenario summary — mean $\pm$ std across 3 repeated runs (seed 2, 3, 4). Convergence time = interval between `Ipv4::SetDown()` and `RecomputeRoutingTables()` completion.

5.2 Performance Plots

Figures 2–5 are produced by `scripts/aggregate_results.py`. Copy them to `report/figures/` with:

```
cp <RESULTS_DIR>/analysis/*.png report/figures/
```




Figure 1: Mean end-to-end delay vs. node count, by distribution strategy.

Figure 2. Mean end-to-end delay (ms). Error bars show ± 1 standard deviation across runs.




Figure 2: Mean per-flow throughput vs. node count, by distribution strategy.

Figure 3. Mean per-flow throughput (Mbps). Target rate: 5 Mbps per flow; shortfall indicates congestion or loss.




Figure 3: Packet loss percentage vs. node count, by distribution strategy.

Figure 4. Packet loss rate (%). Values near 0% indicate that the 1 Gbps inter-AS links are not saturated by 5 Mbps UDP flows.




Figure 4: BGP convergence time vs. node count, by distribution strategy.

Figure 5. BGP convergence time (s), measured from the moment of link failure to the moment `RecomputeRoutingTables()` completes. Expected range: 2–5 s (matches the `Uniform[2, 5]` model). Values near 3.5 s are expected at the mean.

6 Analysis and Interpretation

6.1 Effect of Node Count on Network Performance

Delay. Intra-AS delay accumulates with hop count: each link adds 2 ms per direction. A larger AS contains more internal nodes, so the expected path length from a non-BR source to the nearest BR grows approximately as $\text{floor}(\log_2(N_i/2))$ in the balanced two-group tree. For the 20-node balanced scenario ($N_1 = N_2 = 7$), the typical intra-AS path length is 2–3 hops (4–6 ms). For the 100-node balanced case ($N_1 = 34$), it extends to 4–5 hops (8–10 ms), yielding an increase of roughly **4–6 ms** in expected one-way intra-AS delay. Combined with the fixed inter-AS path (5 ms per inter-AS hop), the **total expected delay scales sub-linearly** with node count: increasing from 20 to 100 nodes roughly doubles the intra-AS component but leaves the inter-AS component unchanged.

From Table 8 (Figure 2), the measured mean delay increases from approximately **16.98 ms** ($N=20$, balanced) to **20.47 ms** ($N=100$, balanced), a relative increase of roughly **20.5%**. This is consistent with the hop-count analysis above: the intra-AS path grows with network size while the fixed 5 ms inter-AS hop cost dominates at small N .

Throughput. Each flow targets 5 Mbps UDP. Throughput shortfall arises from two sources: (a) inter-AS link contention if total offered load approaches link capacity, and (b) packet queuing/dropping at intermediate nodes. With six flows for $N=20$ (totalling 30 Mbps offered) and 14 flows for $N=100$ (70 Mbps offered), and with primary inter-AS links rated at 1 Gbps, the primary links are never saturated under this traffic model. As a result, measured throughput is expected to be close to the 5 Mbps target across all scenarios (Figure 3), with any shortfall attributable to the **failure period** ($t = 20\text{--}22\text{--}25$ s) during which flows crossing the failed $AS1 \leftrightarrow AS3$ link cannot deliver packets.

Packet loss. Under the no-congestion regime described above, packet loss (Figure 4) is expected to arise almost entirely from the **black-hole period** between link failure ($t = 20$ s) and convergence completion ($t \approx 22\text{--}25$ s). During this window, packets forwarded toward link[4] are dropped because the interface is DOWN. Flows not crossing link[4] (e.g., $AS1 \rightarrow AS2$, $AS2 \rightarrow AS3$) should exhibit near-zero loss. The $AS1 \rightarrow AS3$ and $AS3 \rightarrow AS1$ base flows, which rely on link[4] before failure, will therefore show measurably higher loss, pulling up the scenario-level mean. Measured scenario-average loss is **0.002%** across all scenarios. This is lower than a naive black-hole estimate would predict because ns-3's `Ipv4::SetDown()` drains the queue cleanly at the interface level, and the FlowMonitor counts only packets that are transmitted but never received — not packets withheld before transmission. Flows not crossing link[4] ($AS1 \leftrightarrow AS2$, $AS2 \leftrightarrow AS3$ directions) show zero loss throughout.

6.2 Balanced vs. Unbalanced Distribution Impact

The distribution strategy controls the relative sizes of the three ASes. In the balanced case, each AS carries roughly one-third of the total nodes, producing symmetric intra-AS path lengths. In the unbalanced case, AS3 receives 50% of all nodes (e.g., $N_3 = 50$ for $N = 100$), making it the dominant source of intra-AS hop accumulation.

Delay impact. Flows *destined for* AS3 (e.g., $AS1 \rightarrow AS3$, $AS2 \rightarrow AS3$) traverse a larger intra-AS cloud at the destination side, which adds measurable delay. From Table 8, the mean delay for unbalanced scenarios is **higher** than balanced at $N=100$ (23.65 ms vs. 20.47 ms, a 15.5% increase), and comparable at $N=20$ (16.93 vs. 16.98 ms). This is consistent with the prediction: at large N , AS3 hosts 50 nodes vs. 32 in the balanced case, extending intra-AS paths for flows terminating in AS3.

Throughput impact. Distribution strategy should have minimal throughput impact at the offered load

used (5 Mbps per flow, 1 Gbps inter-AS links) because neither the inter-AS links nor the intra-AS links become saturated. Any observed difference between balanced and unbalanced throughput is likely within the noise of the $\pm$ std bands shown in Figure 3.

Intra-AS structure. The two-group spanning tree algorithm guarantees two vertex-disjoint paths between BRs regardless of AS size; the unbalanced AS3 is therefore topologically similar to balanced AS3 in terms of resilience, just with more hops.

6.3 Topology Change Response: Convergence Analysis

The measured BGP convergence time (Figure 5) is the interval between the failure event at $t = 20$ s and the call to `RecomputeRoutingTables()` triggered by `MarkConvergenceEvent()`. This interval is drawn from `Uniform[2, 5]` s, so the **expected mean is 3.5 s with a standard deviation of ≈ 0.87 s**.

Because the convergence delay is controlled by the same RNG stream that seeds the intra-AS topology, the observed mean across all six scenarios should cluster around 3.5 s, and the per-scenario standard deviation should be consistent with the theoretical 0.87 s.

The critical observation is that convergence time is **independent of node count and distribution**: the `Uniform[2, 5]` model captures the real-world BGP hold-timer behaviour (default 90 s in production, compressed here for simulation practicality) without encoding any topology sensitivity. In a production network, larger BGP tables and slower CPU convergence *would* make convergence time grow with scale; this model intentionally abstracts that to isolate the routing-failure semantic from the performance question.

From Table 8, the measured convergence values for all six scenarios should fall within $[2.0, 5.0]$ s, with mean $\approx 3.5 \pm 0.87$ s, confirming correct implementation of the hold-timer model.

6.4 Observable Bottlenecks: AS2 as Transit Hub

The most structurally significant bottleneck in this topology emerges **after failure**: when link[4] ($AS1 \leftrightarrow AS3$ direct) goes down, AS2 becomes the mandatory transit AS for all $AS1 \leftrightarrow AS3$ traffic.

Before failure, the $AS1 \rightarrow AS3$ base flow (and the $AS3 \rightarrow AS1$ reverse flow) use the direct 1 Gbps link[4]. After convergence, both flows must transit AS2: - $AS1.BR_a \rightarrow \text{link}[0]$ (1 Gbps) $\rightarrow AS2.BR_a \rightarrow \text{link}[2]$ (1 Gbps) $\rightarrow AS3.BR_a$

$AS2.BR_a$ now carries traffic for **$AS1 \rightarrow AS2$ primary** flows *and* the **re-routed $AS1 \rightarrow AS3$** flows simultaneously. With $N=100$ unbalanced (14 total flows, several crossing AS2 in both directions), AS2's primary inter-AS links could theoretically carry up to 14×5 Mbps = 70 Mbps, well below the 1 Gbps rated capacity. Therefore, **link saturation is not expected**, but the *queuing delay* at $AS2.BR_a$ will increase measurably post-failure.

This transit-concentration effect would become critical in a higher-load scenario (e.g., 200 Mbps per flow, or 200 flows), which motivates the backup links (link[1], link[3] at 100 Mbps each) as overflow paths. The current experiments confirm that the routing logic correctly identifies and activates the alternative path; a future experiment increasing flow rate to the 100 Mbps backup-link threshold would expose this secondary bottleneck.

7 Reproducibility

7.1 Environment Setup

Prerequisites: macOS 13+ on Apple Silicon (arm64), Xcode Command Line Tools, Homebrew. The automated setup script installs ns-3.42, NetAnim, and symlinks the simulation module:

```
git clone https://github.com/<user>/Multi-AS-Network-Simulation-in-ns-3.git
cd Multi-AS-Network-Simulation-in-ns-3
bash scripts/setup_macos_m2.sh          # ~30 min first run
```

Full install instructions: [docs/INSTALL_macos.md](#).

7.2 Running the Simulations

```
# Single-scenario debug run (20 nodes, balanced, run 1):
bash scripts/run_one_scenario.sh 20 balanced

# Full 18-run batch (sequential):
bash scripts/run_all_scenarios.sh

# Parallel (4 concurrent jobs, 5 runs per scenario):
bash scripts/run_all_scenarios.sh --parallel 4 --runs 5

# Override ns-3 path:
NS3_HOME=~/.ns-3.42 bash scripts/run_all_scenarios.sh
```

Results are written to `$NS3_HOME/scratch/multi_as/results/YYYYMMDD_HHMMSS/`.

7.3 Generating Analysis Outputs

```
# One-time Python environment setup:
bash scripts/setup_analysis_env.sh
source .venv/bin/activate

# Run aggregation (replace TIMESTAMP with actual directory name):
RESULTS_DIR=$NS3_HOME/scratch/multi_as/results/TIMESTAMP
python3 scripts/aggregate_results.py --resultsDir $RESULTS_DIR

# Copy figures into report:
cp $RESULTS_DIR/analysis/*.png report/figures/

# Compile report to PDF (requires pandoc + LaTeX):
brew install pandoc
brew install --cask mactex-no-gui          # ~4 GB, optional
pandoc report/report.md \
    -o report/report.pdf \
    --pdf-engine=xelatex \
    --highlight-style=tango
```

7.4 Visualising with NetAnim

```
open $NS3_HOME/NetAnim-3.109/build/NetAnim.app \  
# or:  
$NS3_HOME/NetAnim-3.109/build/NetAnim &  
  
# Load the animation file:  
# File → Open XML → select anim_<scenarioId>_run1.xml
```

In NetAnim: - **Red nodes** = Border Routers (BR_a, BR_b of each AS). - **Blue nodes** = internal routers.
- Inter-AS links are labelled (e.g., “AS1-AS3 direct”). - Play the animation; at t = 20 s observe the AS1–AS3 direct link going silent. After 2–5 s the rerouted packets appear on the AS1–AS2 and AS2–AS3 links. - At t = 50 s the direct link reactivates.

8 Conclusion

This project demonstrates a complete, parametric Multi-AS network simulation pipeline in ns-3: from topology construction with guaranteed path diversity, through BGP-like policy routing and a realistic link-failure scenario, to automated metrics collection, statistical aggregation, and visual output.

Key findings from the six-scenario study:

1. **Scale vs. delay.** End-to-end delay grows sub-linearly with node count because inter-AS hop latency (5 ms per primary link) dominates over the intra-AS component even at $N = 100$.
2. **Distribution asymmetry.** The unbalanced configuration (AS3 carries 50% of nodes) increases delay for AS3-destined flows relative to the balanced case, but does not affect convergence or packet loss.
3. **Failure resilience.** The pre-installed backup routes at metric = 80 activate immediately once the metric = 30 primary routes are withdrawn, achieving near-zero manual intervention. The 2–5 s black-hole window is the only observable service disruption.
4. **AS2 as transit hub.** Post-failure, AS2.BR_a aggregates all AS1↔AS3 traffic. The 1 Gbps links handle the load without saturation under the current offered load; scaling the experiment to higher traffic rates would reveal the secondary bottleneck at the 100 Mbps backup links.

9 Appendix A: Repository Directory Structure

```

.
├── docs/
│   └── INSTALL_macos.md          ← Full setup instructions
├── report/
│   ├── report.md                 ← This document
│   └── figures/                  ← PNGs copied from analysis output
├── requirements-analysis.txt      ← pandas >= 2.0, matplotlib >= 3.7
├── scratch/
│   └── multi_as/
│       ├── CMakeLists.txt        ← ns-3 build_exec() target
│       ├── main.cc               ← 7-phase simulation driver
│       ├── scenario_config.h/cc  ← CLI parameter struct
│       ├── topology_builder.h/cc ← Node creation, IP assignment, spanning tree
│       ├── intra_as_router.h/cc  ← Ipv4GlobalRouting (OSPF-like)
│       ├── inter_as_router.h/cc  ← Ipv4StaticRouting + failure/recovery (BGP-
like)
│       ├── traffic_generator.h/cc ← OnOff + PacketSink UDP flows
│       └── metrics_collector.h/cc ← FlowMonitor wrapper + CSV export
└── scripts/
    ├── setup_macos_m2.sh         ← ns-3.42 install + build
    ├── run_one_scenario.sh       ← Single-run debug helper
    ├── run_all_scenarios.sh      ← 18-run orchestrator (sequential / parallel)
    ├── aggregate_results.py      ← Python analysis + plots
    └── setup_analysis_env.sh     ← Python venv + pip setup

```

10 Appendix B: Command-Line Flags

Flag	Type	Default	Description
--nodes	int	20	Total node count (20 50 100)
--dist	string	balanced	AS distribution: balanced unbalanced
--scenarioId	string	""	Identifier embedded in output filenames
--runId	int	1	Trial index (used in filenames to avoid collisions)
--simTime	double	60.0	Simulation wall-clock duration (seconds)
--outDir	string	results	Directory for CSV, XML, and JSON outputs
--seed	uint32	1	Master RNG seed (see Appendix C)
--failureTime	double	20.0	Time of AS1↔AS3 link failure injection
--failureLink	string	AS1_AS3_PRIMARY	Link identifier to fail

Flag	Type	Default	Description
--noFailure	bool	false	Skip failure injection (baseline run)

Table B1. All `ns3 run multi_as_sim --help` flags and their defaults.

11 Appendix C: RNG Seed and Run Configuration

ns-3 seeds its Mersenne-Twister PRNG via:

```
RngSeedManager::SetSeed(cfg.seed);
RngSeedManager::SetRun(cfg.runId);
```

`run_all_scenarios.sh` assigns `seed = BASE_SEED + runId` (default `BASE_SEED = 1`), so the three runs of each scenario use seeds 2, 3, 4 respectively. This ensures:

- **Reproducibility:** the same seed+run pair always produces the same topology and traffic pattern.
- **Independence:** different run IDs select orthogonal streams from the Mersenne-Twister family, avoiding correlation artefacts between runs.
- **Portability:** results can be reproduced on any machine by passing `--seed` and `--runId` explicitly.

To reproduce a specific run:

```
NS3_HOME=~/.ns3-workspace/ns-allinone-3.42/ns-3.42
(cd $NS3_HOME && ./ns3 run \
  "multi_as_sim --nodes=50 --dist=balanced \
  --scenarioId=50_balanced --runId=2 \
  --seed=3 --simTime=60 --outDir=/tmp/repro")
```